

Computer Science CS134C (Spring 2018)

Duane A. Bailey

Laboratory 6

Building an Oracle (due 11pm, Wednesday)

Objective. To build a simple class that generates text in an intelligent (?) manner.

This week we'll write a simple little class, an `Oracle`, that can be trained to generate "readable" random text. The class makes use of a technique that *fingerprints* a corpus by keeping track of a distribution of combinations of n letters.

The Tasks. Here are the steps to completing this week's lab.

1. Download the starter kit for this package in the usual manner:

```
git clone ssh://18xyz@davey.cs.williams.edu/~dab/18xyz/lab6.git lab6
```

This kit is fairly minimal; you'll be writing most of the code from scratch.

2. In a file `oracle.py` write a class, `Oracle`, that inherits from the `object` class. It would typically be used in the following manner:

```
o = Oracle(window=3)
text = ' '.join([line.strip() for line in open('tomsawyer.txt')])
o.intern(text)
print(o.generate('Tom'))
```

The initializer for the `Oracle` takes a window size, w . When the `Oracle` scans a corpus of text, it scans the text from beginning to end keeping track of the frequency of each combination of w characters. We'll think of this, essentially, as a *fingerprint* of the corpus.

The `generate` method, given a seed of at least $w-1$ characters, generates a string of text, randomly, given the fingerprint internalized in the `Oracle`.

3. Write the `__init__` method. This method should accept a window size that defaults to 3. The distribution for the `Oracle` is initially uniform; there's no real fingerprint.
4. Write a method, `o.intern(s)`, that takes a string, s , and records the occurrences of w character combinations, where w is the window size of the `Oracle`.
5. Write a method, `o.follows(s)`. The string should be at least $w-1$ characters long, where w is the `Oracle`'s window size. It returns a random character that, given knowledge of the interned text, is likely to follow s . For example, if the window size is 3, then

```
>>> o = Oracle()
>>> o.intern('this, that, and the other thing')
>>> o.follows('th')
'e'
```

another possibilities would be 'i' or 'a'. You should think about what happens if the seed string does not appear in the interned distribution.

6. Write a method, `o.generate(seed)`, that, given a string of at least length `w-1` generates a string that follows the distribution of the original corpus.

The length of the string should be bounded by a default `length`. For example, `length` might be 80 if you're generating pages of lines of text. The `length` might be 140 if you wanted to use this to generate tweets.

7. To demonstrate the functionality of your Oracle, write code (guarded by an `if __name__` statement), that prints a page of text motivated by the fingerprint of Twain's *Tom Sawyer*. Or you may use another text (I've included Jane Austen's *Pride and Prejudice*, if you're so inclined).
8. Feel free to add other features to the Oracle. Are there properties that might be interesting? How might you simplify the interning of text from a file like `tomsawyer.txt`?
9. Make sure that you provide appropriate documentation strings for the class and its methods.

Remember: use git to add and commit changes to files you add to this project.

★